

ANALYTICAL SQL PROJECT

[Analyzing Store Transactions Database]



Abdullah Attia

Table of content:

1- EDA

2- RFM

3- some insights

1- Exploratory Data Analysis

Let's do some basic EDA!

First, Getting to know the data with some data discovery queries.

```
SELECT
  *
FROM
  TABLERETAIL
LIMIT
  5;
```

| Invoice character varying (50) 🔒 | stockcode character varying (50) 🔒 | quantity integer 🔒 | Invoicedate character varying (50) 🔒 | price double precision 🔒 | customer_id character varying (50) 🔒 |
|-------------------------------------|---------------------------------------|-----------------------|---|-----------------------------|---|
| 537215 | 85124C | 12 | 12/5/2010 15:38 | 2.55 | 12747 |
| 537215 | 85124B | 6 | 12/5/2010 15:38 | 2.55 | 12747 |
| 537215 | 84879 | 16 | 12/5/2010 15:38 | 1.69 | 12747 |
| 537215 | 85062 | 24 | 12/5/2010 15:38 | 1.65 | 12747 |
| 537215 | 85064 | 6 | 12/5/2010 15:38 | 5.45 | 12747 |

So we have orders dataset of transactions, with order item as level of granularity,

```

SELECT
    COUNT(DISTINCT INVOICE) INVOICE_COUNT,
    COUNT(DISTINCT STOCKCODE) STOCKCODE_COUNT,
    COUNT(DISTINCT CUSTOMER_ID) CUSTOMER_COUNT,
    MIN(INVOICEDATE::DATE) AS "first order",
    MAX(INVOICEDATE::DATE) AS "last order",
    MAX(INVOICEDATE::DATE) - MIN(INVOICEDATE::DATE) "dataset time period"
FROM
    TABLERETAIL;

```

| invoice_count bigint | stockcode_count bigint | customer_count bigint | first order date | last order date | dataset time period integer |
|-------------------------|---------------------------|--------------------------|---------------------|--------------------|--------------------------------|
| 717 | 2335 | 110 | 2010-12-01 | 2011-12-09 | 373 |

So in this dataset,

We have 717 orders including 2335 different products made by 110 different customers over around a year.

Now, Let's Dive a little bit in the data to get some more insights and answer some questions!

1- what are the top selling products

```
SELECT
    STOCKCODE,
    SUM(QUANTITY) "UNITS SOLD"
FROM
    TABLERETAILE T
GROUP BY
    STOCKCODE
ORDER BY
    "UNITS SOLD" DESC
LIMIT
    5;
```

| stockcode character varying (50) 🔒 | UNITS SOLD bigint 🔒 |
|---------------------------------------|------------------------|
| 84077 | 7824 |
| 84879 | 6117 |
| 22197 | 5918 |
| 21787 | 5075 |
| 21977 | 4691 |

So these are the top selling Products, we can do some further analysis, to find out why?
Is it related to their prices, quality, promotions, and more.

2- Best Selling Month

a. In terms of Number of Orders

```
SELECT
    TO_CHAR(INVOICEDATE::DATE, 'Mon') "month",
    COUNT(INVOICE) "Order Count"
FROM
    TABLERETAIL T
GROUP BY
    "month"
ORDER BY
    "Order Count" DESC;
```

| month text | Order Count bigint |
|---------------|-----------------------|
| Nov | 3200 |
| Dec | 1955 |
| Sep | 1411 |
| Oct | 1003 |
| May | 983 |
| Jun | 745 |
| Jul | 743 |
| Mar | 727 |
| Aug | 595 |
| Feb | 556 |
| Apr | 479 |
| Jan | 461 |

b. In terms of Sales

```
SELECT
    TO_CHAR(INVOICEDATE::DATE, 'Mon') "month",
    ROUND(SUM(QUANTITY * PRICE)::NUMERIC, 2) AS SALES
FROM
    TABLERETAILE T
GROUP BY
    "month"
ORDER BY
    SALES DESC;
```

| month | sales |
|-------|----------|
| text | numeric |
| Nov | 45633.38 |
| Aug | 38374.64 |
| Sep | 27853.82 |
| Dec | 24547.09 |
| Oct | 19735.07 |
| May | 19496.18 |
| Mar | 17038.01 |
| Jul | 15664.54 |
| Jun | 13517.01 |
| Feb | 13336.84 |
| Apr | 10980.51 |
| Jan | 9541.29 |

As we can see November was the Best Month In terms of Both Sales And Orders Traffic, however the second and third are different in terms of number of orders, and sales meaning that the average order value for months is changing by a good amount.

3- Best Selling Day

a. In terms of Number of Orders

```
SELECT
    INVOICEDATE::DATE "date",
    COUNT(INVOICE) "Order Count"
FROM
    TABLERETAIL T
GROUP BY
    "date"
ORDER BY
    "Order Count" DESC
LIMIT
    5;
```

| date | Order Count |
|------------|-------------|
| date | bigint |
| 2011-11-23 | 426 |
| 2011-11-17 | 304 |
| 2011-09-11 | 289 |
| 2011-11-25 | 268 |
| 2011-11-15 | 268 |

b. In terms of Sales

```
SELECT
    INVOICEDATE::DATE "day",
    ROUND(SUM(QUANTITY * PRICE)::NUMERIC, 2) AS SALES
FROM
    TABLERETAIL T
GROUP BY
    "day"
ORDER BY
    SALES DESC
LIMIT
    5;
```

| day | sales |
|------------|----------|
| date | numeric |
| 2011-08-04 | 19304.68 |
| 2011-08-11 | 9349.72 |
| 2011-11-09 | 6676.09 |
| 2011-09-22 | 4839.38 |
| 2011-11-17 | 4672.72 |

factors such as promotions, marketing campaigns, special events, and customer behavior can influence these differences. Analyzing both metrics separately can provide insights into different aspects of daily sales performance and help in making informed business decisions.

4- Traffic Over the Hour Of Day

a. Number of Orders

```
SELECT
    TO_CHAR(INVOICEDATE::TIMESTAMP, 'HH24') AS "Hour",
    COUNT(INVOICE) AS "count of orders"
FROM
    TABLERETAIL
GROUP BY
    TO_CHAR(INVOICEDATE::TIMESTAMP, 'HH24')
ORDER BY
    "Hour";
```

| Hour text | count of orders bigint |
|--------------|---------------------------|
| 07 | 1 |
| 08 | 80 |
| 09 | 550 |
| 10 | 944 |
| 11 | 1713 |
| 12 | 2439 |
| 13 | 2173 |
| 14 | 1710 |
| 15 | 1827 |
| 16 | 698 |
| 17 | 372 |
| 18 | 120 |
| 19 | 199 |
| 20 | 32 |

b. Sales

```
SELECT
    TO_CHAR(INVOICEDATE::TIMESTAMP, 'HH24') AS "Hour",
    ROUND(SUM(QUANTITY * PRICE)::NUMERIC, 2) AS SALES
FROM
    TABLERETAIL
GROUP BY
    TO_CHAR(INVOICEDATE::TIMESTAMP, 'HH24')
ORDER BY
    "Hour"
```

| Hour text | sales numeric |
|--------------|------------------|
| 07 | 535.50 |
| 08 | 2065.29 |
| 09 | 15988.75 |
| 10 | 28699.67 |
| 11 | 21820.18 |
| 12 | 39473.63 |
| 13 | 44414.87 |
| 14 | 24235.95 |
| 15 | 35772.63 |
| 16 | 11249.71 |
| 17 | 8260.93 |
| 18 | 21755.01 |
| 19 | 1362.57 |
| 20 | 83.69 |

The sales amount is highest during the late morning and early afternoon hours, specifically between 10 and 14. This suggests that these hours are the busiest in terms of sales activity

However, after this time sales start to decrease so we might want to figure out why or add some promotions and discounts to try to push the sales further.

5- Top 5 Customers

a. In terms of Number of Orders

```
SELECT
  CUSTOMER_ID,
  COUNT(*) "count of orders"
FROM
  TABLERETAIL T
GROUP BY
  CUSTOMER_ID
ORDER BY
  "count of orders" DESC
LIMIT
  5;
```

| customer_id character varying (50) 🔒 | count of orders bigint 🔒 |
|---|-----------------------------|
| 12748 | 4596 |
| 12921 | 720 |
| 12867 | 538 |
| 12841 | 420 |
| 12856 | 314 |

b. In terms of Sales

```
SELECT
  CUSTOMER_ID,
  ROUND(SUM(quantity*price)::numeric, 2) sales
FROM
  TABLERETAIL T
GROUP BY
  CUSTOMER_ID
ORDER BY
  sales DESC
LIMIT
  5;
```

| customer_id character varying (50) 🔒 | sales numeric 🔒 |
|---|--------------------|
| 12931 | 42055.96 |
| 12748 | 33719.73 |
| 12901 | 17654.54 |
| 12921 | 16587.09 |
| 12939 | 11581.80 |

We can spot the top customers through some customer Segmentation or even RFM analysis to identify their behavior and give them some sort of discounts to motivate their behavior more and see their percentage from the total income, etc. Will see more on this later.

2- RFM

And now we will segment all the customers in the database based on their behavior, mainly focusing on:

- Recency:
When was the most recent order?
- Frequency:
How often does the customer buy from the Business?
- Monetary:
How much has the customer paid to the Business in total?

Normally for each Value we have a Score, then we combine these into some groups that we add the customers to.

However, in this Analysis, we will combine the Frequency and Monetary scores as they sort of decide the value of the customer (the customer importance)

And the recency as it is what decides how recent the customer and its status (recent, lost, etc.)

So, Based on these 2 values for each customer.

We will decide the customer segment.

| Group name | Recency score | AVG(Frequency & Monetary) score |
|-----------------------------|---------------|----------------------------------|
| Champions | 5 | 5 |
| | 5 | 4 |
| | 4 | 5 |
| Potential Loyalists | 5 | 2 |
| | 4 | 2 |
| | 3 | 3 |
| | 4 | 3 |
| Loyal Customers | 5 | 3 |
| | 4 | 4 |
| | 3 | 5 |
| | 3 | 4 |
| Recent Customers | 5 | 1 |
| Promising | 4 | 1 |
| | 3 | 1 |
| Customers Needing Attention | 3 | 2 |
| | 2 | 3 |
| | 2 | 2 |
| At Risk | 2 | 5 |
| | 2 | 4 |
| | 1 | 3 |
| Cant Lose Them | 1 | 5 |
| | 1 | 4 |
| Hibernating | 1 | 2 |
| Lost | 1 | 1 |

Now the Code:

```
WITH
RFM_STATS AS (
  SELECT
    CUSTOMER_ID AS CUSTOMER,
    MAX(INVOICEDATE::DATE) AS "latest_order",
    CURRENT_DATE - MAX(INVOICEDATE::DATE) AS RECENCY,
    COUNT(DISTINCT INVOICE) AS FREQUENCY,
    ROUND(SUM(PRICE * QUANTITY)::NUMERIC, 2) AS MONETARY
  FROM
    TABLERETAILE T
  GROUP BY
    CUSTOMER
),
```

| customer character varying (50) | latest_order date | recency integer | frequency bigint | monetary numeric |
|------------------------------------|----------------------|--------------------|---------------------|---------------------|
| 12747 | 2011-12-07 | 4472 | 11 | 4196.01 |
| 12748 | 2011-12-09 | 4470 | 210 | 33719.73 |
| 12749 | 2011-12-06 | 4473 | 5 | 4090.88 |

This CTE just gets the customer and its status that will be scored on later.

```

RFM_SCORES AS (
  SELECT
    CUSTOMER,
    "latest_order",
    RECENCY,
    NTILE(5) OVER (
      ORDER BY
        RECENCY DESC
    ) AS R_SCORE,
    FREQUENCY,
    NTILE(5) OVER (
      ORDER BY
        FREQUENCY
    ) AS F_SCORE,
    MONETARY,
    ROUND(
      PERCENT_RANK() OVER (
        ORDER BY
          MONETARY
      )::NUMERIC,
      3
    ) AS M_SCORE
  FROM
    RFM_STATS
),

```

| customer character varying (50) | latest_order date | recency integer | r_score integer | frequency bigint | f_score integer | monetary numeric | m_score numeric |
|------------------------------------|----------------------|--------------------|--------------------|---------------------|--------------------|---------------------|--------------------|
| 12855 | 2010-12-02 | 4842 | 1 | 1 | 1 | 38.10 | 0.000 |
| 12967 | 2010-12-16 | 4828 | 1 | 2 | 2 | 1660.90 | 0.706 |
| 12829 | 2011-01-07 | 4806 | 1 | 2 | 3 | 293.00 | 0.156 |

Setting the Scores based on the customer stats.

```

R_FM_SCORES AS (
  SELECT
    CUSTOMER,
    REGENCY,
    FREQUENCY,
    MONETARY,
    M_SCORE,
    R_SCORE,
    ROUND((SUM(M_SCORE + F_SCORE) / 2)::NUMERIC) AS FM_SCORE
  FROM
    RFM_SCORES
  GROUP BY
    CUSTOMER,
    REGENCY,
    FREQUENCY,
    MONETARY,
    M_SCORE,
    R_SCORE
)

```

| customer character varying (50) | recency integer | frequency bigint | monetary numeric | m_score numeric | r_score integer | fm_score numeric |
|------------------------------------|--------------------|---------------------|---------------------|--------------------|--------------------|---------------------|
| 12747 | 4472 | 11 | 4196.01 | 0.917 | 5 | 3 |
| 12748 | 4470 | 210 | 33719.73 | 0.991 | 5 | 3 |
| 12749 | 4473 | 5 | 4090.88 | 0.899 | 5 | 2 |

Combining the FREQUENCY and MONETARY Values


```

SELECT
  CUSTOMER,
  REGENCY,
  FREQUENCY,
  MONETARY,
  R_SCORE,
  FM_SCORE,
  CASE
    WHEN R_SCORE = 5
      AND FM_SCORE IN (5, 4) THEN 'champions'
    WHEN R_SCORE = 4
      AND FM_SCORE = 5 THEN 'champions'
    WHEN R_SCORE = 5
      AND FM_SCORE = 2 THEN 'potential loyalist'
    WHEN R_SCORE = 4
      AND FM_SCORE IN (2, 3) THEN 'potential loyalist'
    WHEN R_SCORE = 3
      AND FM_SCORE = 3 THEN 'potential loyalist'
    WHEN R_SCORE = 5
      AND FM_SCORE = 3 THEN 'loyal customers'
    WHEN R_SCORE = 4
      AND FM_SCORE = 4 THEN 'loyal customers'
    WHEN R_SCORE = 3
      AND FM_SCORE IN (5, 4) THEN 'loyal customers'
    WHEN R_SCORE = 5
      AND FM_SCORE = 1 THEN 'recent customer'
    WHEN R_SCORE = 4
      AND FM_SCORE = 1 THEN 'promising'
    WHEN R_SCORE = 3
      AND FM_SCORE = 1 THEN 'promising'
    WHEN R_SCORE = 2
      AND FM_SCORE IN (3, 2) THEN 'needs attention'
    WHEN R_SCORE = 3
      AND FM_SCORE = 2 THEN 'needs attention'
    WHEN R_SCORE = 2
      AND FM_SCORE IN (5, 4) THEN 'At Risk'
    WHEN R_SCORE = 1
      AND FM_SCORE = 3 THEN 'At Risk'
    WHEN R_SCORE = 1
      AND FM_SCORE IN (5, 4) THEN 'cant lose them'
    WHEN R_SCORE = 1
      AND FM_SCORE = 2 THEN 'Hibernating'
    WHEN R_SCORE = 2
      AND FM_SCORE = 1 THEN 'Hibernating'
    WHEN R_SCORE = 1
      AND FM_SCORE = 1 THEN 'lost'
  END AS CUSTOMER_SEGMENTATION
FROM
  R_FM_SCORES
ORDER BY
  CUSTOMER_SEGMENTATION;

```

| customer character varying (50) | recency integer | frequency bigint | monetary numeric | r_score integer | fm_score numeric | customer_segmentation text |
|------------------------------------|--------------------|---------------------|---------------------|--------------------|---------------------|-------------------------------|
| 12830 | 4508 | 6 | 6814.64 | 3 | 2 | needs attention |
| 12822 | 4541 | 2 | 948.88 | 3 | 2 | needs attention |
| 12856 | 4478 | 6 | 2179.93 | 5 | 2 | potential loyalist |

And now the final Segmentation.

3- Answering Business Questions

- First Question)

What is the Maximum Number of Consecutive Days a customer made Purchases?

```
WITH
  T1 AS (
    SELECT
      CUSTOMER,
      AMOUNT,
      ORDER_DATE,
      EXTRACT(
        DAY
        FROM
          (
            ORDER_DATE - (
              ROW_NUMBER() OVER (
                PARTITION BY
                  CUSTOMER
                ORDER BY
                  ORDER_DATE
              ) * INTERVAL '1 day'
            )
          )
      ) AS DIFFERENCE
    FROM
      TRANSACTIONS
  ),
  MAX_CON_DAYS AS (
    SELECT
      CUSTOMER,
      ORDER_DATE,
      COUNT(DIFFERENCE) OVER (
        PARTITION BY
          DIFFERENCE,
          CUSTOMER
      ) CONSECUTIVE_DAYS
    FROM
      T1
    WHERE
      AMOUNT > 0
  )
SELECT
  CUSTOMER,
  MAX(CONSECUTIVE_DAYS) MAX_CON_DAYS
FROM
  MAX_CON_DAYS
GROUP BY
  CUSTOMER
ORDER BY
  MAX_CON_DAYS DESC;
```

| customer integer | max_con_days bigint |
|---------------------|------------------------|
| 105358175 | 61 |
| 170731910 | 61 |
| 13503299 | 60 |

The Algorithm Behind:

First, we are getting a base to the days that are calculating on HOW?

Let's say we're having a customer who has the following Pattern.

1/1
2/1
3/1
5/1
6/1

We can add a row number first

1/1 - 1
2/1 - 2
3/1 - 3
5/1 - 4
6/1 - 5

Then subtract the date from the row number as difference:

1/1 - 1, 0
2/1 - 2, 0
3/1 - 3, 0
5/1 - 4, 1
6/1 - 5, 1

And this way we are grouping the consecutive days, and customer for using partition by and counting each group.

Then Removing the people who didn't purchase.

- Second Question)

On average, how many transactions does it take a customer to reach a spent threshold of 250 L.E?

```

WITH
  RUNNING_TOTAL AS (
    SELECT
      CUSTOMER,
      SUM(AMOUNT) OVER (
        PARTITION BY
          CUSTOMER
        ORDER BY
          ORDER_DATE
      ) RUNNING_TOTAL_SPENT,
      ORDER_DATE
    FROM
      TRANSACTIONS T
  ),
  TARGET_CUSTOMERS AS ( --customers who spent more than 250
    SELECT DISTINCT
      CUSTOMER
    FROM
      RUNNING_TOTAL
    WHERE
      RUNNING_TOTAL_SPENT > 250
  ),
  DAYS_BEFORE_250 AS (
    SELECT
      CUSTOMER,
      RUNNING_TOTAL_SPENT,
      ROW_NUMBER() OVER (
        PARTITION BY
          CUSTOMER
        ORDER BY
          ORDER_DATE
      ) AS ORDER_COUNT_BEFORE_250
    FROM
      RUNNING_TOTAL
    WHERE
      RUNNING_TOTAL_SPENT < 250
      AND CUSTOMER IN (
        SELECT
          CUSTOMER
        FROM
          TARGET_CUSTOMERS
      )
  ),
  AVG_DAYS_BEFORE_250 AS (
    SELECT
      CUSTOMER,
      MAX(ORDER_COUNT_BEFORE_250) AS NUMBER_OF_ORDERS_BEFORE_250
    FROM
      DAYS_BEFORE_250
    GROUP BY
      CUSTOMER
  )
SELECT
  ROUND(AVG(NUMBER_OF_ORDERS_BEFORE_250)::NUMERIC)
FROM
  AVG_DAYS_BEFORE_250

```

The Algorithm Behind:

We first calculate the running total per transaction for each customer. Then, we isolate the running total of transactions before reaching the 250 threshold per customer. For customers who surpass this threshold, we assign row numbers to transactions and identify the maximum number within each group. Finally, we compute the average of these maximum numbers across all customers, providing insights into transaction behavior around the 250 thresholds.